

THREE FORWARD, ONE BACKWARD: MEMORY-EFFICIENT FULL-RANK FINE-TUNING OF LARGE MODELS VIA EXTRA FORWARD PASSES

Anonymous authors

Paper under double-blind review

ABSTRACT

Fine-tuning large language models (LLMs) has achieved significant success in downstream tasks. However, as the model size continues to grow, traditional fine-tuning methods have become increasingly impractical due to their high computational and memory costs. This has motivated researchers to explore parameter-efficient and memory-friendly fine-tuning strategies to enable scalable approaches, with Low-Rank Adaptation (LoRA) standing out as a representative work. However, the LoRA update is restricted to a low-rank subspace, which results in suboptimal performance compared to the full-parameter update. Recent research has also explored memory-efficient fine-tuning LLMs using just forward passes while suffer from high variance in gradient estimation and low convergence speed. To address the issues above, we propose a new alternating optimization framework called LMAO (Low-rank and Memory-efficient Zeroth-Order Alternating Optimization), which combines the advantages of LoRA and MeZO. This method alternately updates the low-rank components and zeroth-order directions during training. By performing three forward propagations and one backward propagation, each update is full-rank, thereby reducing feature loss and enabling efficient fine-tuning under strict memory constraints. We provide theoretical guarantees on the convergence and convergence rate of this method. Empirical results demonstrate that, in experiments on multiple models (e.g., OPT, RoBERTa-large), LMAO achieves performance comparable to first-order methods. This presents a practical and scalable solution for fine-tuning large-scale models. Our source code is available at <https://anonymous.4open.science/r/mlao-C2EC/>.

1 INTRODUCTION

Large language models (LLMs) have achieved remarkable performance across a wide range of domains (Solaiman et al., 2019; Brown et al., 2020; Achiam et al., 2023). However, as LLMs continue to scale, gradient computation can demand over $12\times$ the memory required for inference (Malladi et al., 2023), posing significant challenges to model training and fine-tuning — particularly exacerbating optimization difficulty in resource-constrained environments. This challenge has spurred the development of memory-efficient alternatives, such as parameter-efficient fine-tuning (PEFT), which aim to reduce resource requirements while preserving task-specific performance.

PEFT methods allow the adaptation of LLMs to downstream tasks by tuning only a subset of model parameters or by adding task-specific modules, such as adapters (Houlsby et al., 2019), prefix embeddings (Li & Liang, 2021), and prompt tuning (Lester et al., 2021), without altering the base architecture. Studies further demonstrate that fine-tuning pretrained models can be effectively conducted within reduced-dimensional parameter subspaces (Aghajanyan et al., 2020). Low-Rank Adaptation (LoRA), a prominent PEFT technique, leverages this insight by performing adaptations within optimized low-dimensional parameter spaces. This strategy significantly enhances the efficiency of adapting large-scale pretrained models, validating the theoretical foundations regarding low-rank representations in neural networks (Hu et al., 2022; Aghajanyan et al., 2020; Li et al., 2018). Despite the efficiency, LoRA still falls short of matching full-model fine-tuning performance due to its limited representation power of fine-tuned language models (Hu et al., 2022).

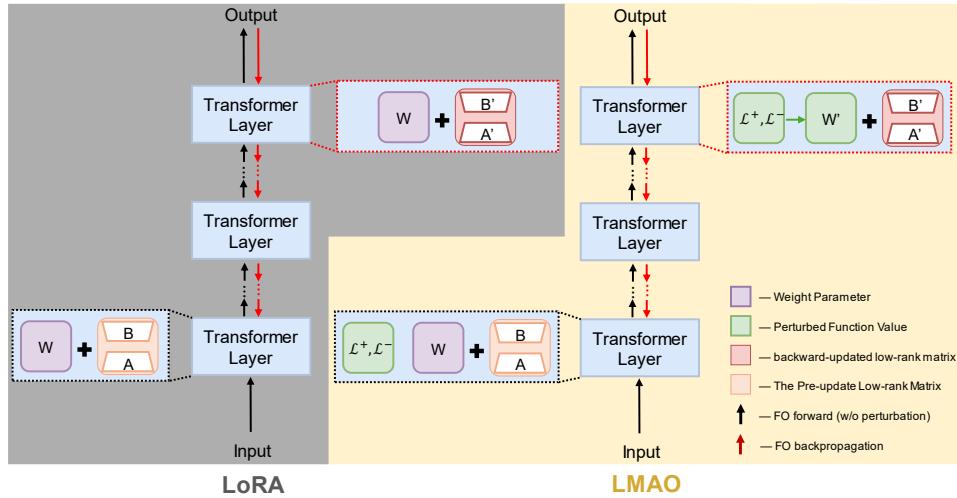


Figure 1: Comparison of standard LoRA and LMAO: LoRA performs a single forward-backward pass, while LMAO alternates it with two zeroth-order forward passes. Note: the forward and backward passes update only the low-rank matrices B and A in LoRA, while the base weight matrix W is updated solely through the forward pass.

Meanwhile, zeroth-order optimization offers an alternative approach to model adaptation, particularly when gradient access is restricted. Zeroth-order optimizers approximate gradients through forward passes alone, eliminating the need for backpropagation and further reducing memory requirements. MeZO (Malladi et al., 2023), a Memory-efficient Zeroth-Order optimizer, established the feasibility of fine-tuning large models using forward passes only, enabling gradient-free adaptation for various downstream tasks. However, the inherent noise and high variance of zeroth-order gradient estimates pose limitations on convergence speed and task performance.

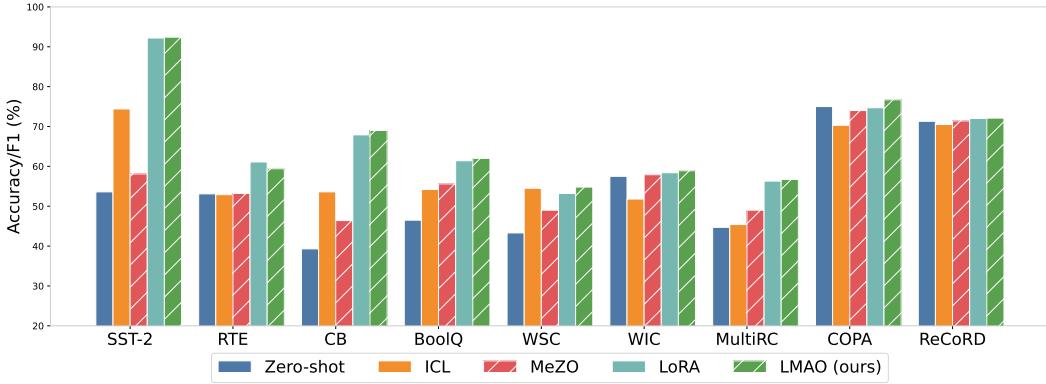


Figure 2: OPT-1.3B results under zero-shot, ICL, MeZO (Memory-efficient Zeroth-Order optimization), LoRA (Low-Rank Adaptation), and LMAO (ours). LMAO achieves superior performance on most tasks. See Table 2 for details.

In this paper, we propose a memory-efficient full-rank fine-tuning method, the Low-rank and Memory-efficient Zeroth-order Alternating Optimization (LMAO) algorithm. LMAO combines LoRA and MeZO to optimize model performance by alternately updating low-rank matrix components and zeroth-order directions. Each iteration involves three forward passes and one backpropagation, ensuring full-rank updates. This approach reduces feature loss and enhances model expressiveness compared to traditional low-rank methods. LMAO also integrates the memory efficiency of both PEFT methods, ensuring tight control over memory usage. It introduces a new alternating optimization framework with theoretical convergence guarantees, making it a memory-efficient solution for fine-tuning large models. The algorithm is well-suited for memory-constrained environments,

108 maintaining high performance while staying within memory limits, enabling efficient large-scale
 109 model optimization.
 110

111 2 RELATED WORK 112

113 **Parameter-Efficient Fine-Tuning.** Parameter-Efficient Fine-Tuning (PEFT) methods reduce the
 114 computational and memory costs of fine-tuning large pre-trained language models (LLMs). PEFT
 115 adapts LLMs to downstream tasks by adjusting a small subset of the model’s parameters while
 116 keeping most weights fixed. This approach achieves task-specific improvements without the extensive
 117 hardware required for full-model fine-tuning. Prompt tuning (Lester et al., 2021) optimizes continuous
 118 prompt vectors appended to input embeddings, enabling task-specific conditioning. Prefix tuning
 119 (Li & Liang, 2021) introduces trainable prefix tokens at each transformer layer for lightweight
 120 conditioning. Adapters (Houlsby et al., 2019) add small neural modules between transformer layers,
 121 which are tuned independently. Low-Rank Adaptation (LoRA) (Hu et al., 2022) reduces parameter
 122 overhead by injecting low-rank matrices into attention and feedforward layers.

123 **Zeroth-Order Fine-Tuning of LLMs.** Zeroth-order (ZO) optimization methods enable gradient-free
 124 fine-tuning, ideal for scenarios where gradient access is infeasible or computationally expensive.
 125 MeZO (Malladi et al., 2023), a Memory-Efficient Zeroth-Order optimizer, shows that ZO fine-tuning
 126 can achieve memory efficiency similar to model inference with forward-pass-only updates. MeZO
 127 uses Simultaneous Perturbation Stochastic Approximation (SPSA) (Spall, 1992) for gradient estimation,
 128 allowing updates without backpropagation. However, MeZO faces challenges in convergence
 129 speed due to high variance in ZO gradient estimates. Recent improvements to MeZO focus on sparsity
 130 (Liu et al., 2024; Guo et al., 2024), variance reduction (Gautam et al., 2024), low-rank structures (Yu
 131 et al., 2024b; Chen et al., 2024), and Hessian information (Zhao et al., 2024; Yu et al., 2024a).

132 3 METHODOLOGY 133

134 3.1 PRELIMINARIES 135

136 **LoRA.** Low-Rank Adaptation of Large Language Models (Hu et al., 2022) is a classical fine-tuning
 137 method for large language models. By introducing trainable low-rank matrices, LoRA reduces the
 138 number of trainable parameters to less than 0.1% of the full model, decreases GPU memory usage by
 139 two-thirds, and introduces no additional inference latency. The study by Aghajanyan et al. (2020)
 140 demonstrates that pretrained models possess an extremely low intrinsic dimensionality, indicating that
 141 fine-tuning a small number of parameters in a low-dimensional subspace can achieve performance
 142 comparable to full-parameter fine-tuning. The core idea is that the parameter update ΔW with respect
 143 to the pretrained weight matrix $W_0 \in \mathbb{R}^{m \times n}$ can be represented via a low-rank decomposition, i.e.,
 144

$$W_0 + \Delta W = W_0 + \frac{\alpha}{r} BA, \quad B \in \mathbb{R}^{m \times r}, \quad A \in \mathbb{R}^{r \times n} \text{ and } r \ll \min(m, n).$$

145 **MeZO.** Although traditional ZO methods can estimate gradients using only two forward passes,
 146 they suffer from the burden of additional memory consumption induced by random perturbations,
 147 which have the same shape as the model parameters. To address this limitation, Memory-efficient
 148 Zeroth-Order optimizer (MeZO) (Malladi et al., 2023) utilizes the random seed trick to generate
 149 random perturbations sequentially on-the-fly, enabling in-place optimization without modifying the
 150 model architecture and achieving memory usage equivalent to inference during fine-tuning. Moreover,
 151 MeZO exhibits broad compatibility, supporting both full-parameter fine-tuning and parameter-efficient
 152 techniques such as LoRA.
 153

154 3.2 LOW-RANK AND MEMORY-EFFICIENT ZEROTH-ORDER ALTERNATING OPTIMIZATION 155

156 Although LoRA reduces trainable parameters by restricting updates to a low-dimensional subspace,
 157 this constraint limits the representational capacity of fine-tuned models, often yielding suboptimal
 158 performance compared to full-parameter updates (Hu et al., 2022). In contrast, MeZO avoids gradient
 159 computations but suffers from high-variance gradient estimates, leading to slower convergence (Yu
 160 et al., 2024a). To leverage the strengths of both, we propose the Low-rank and Memory-efficient
 161 zeroth-order Alternating Optimization (LMAO) framework, which strategically alternates between

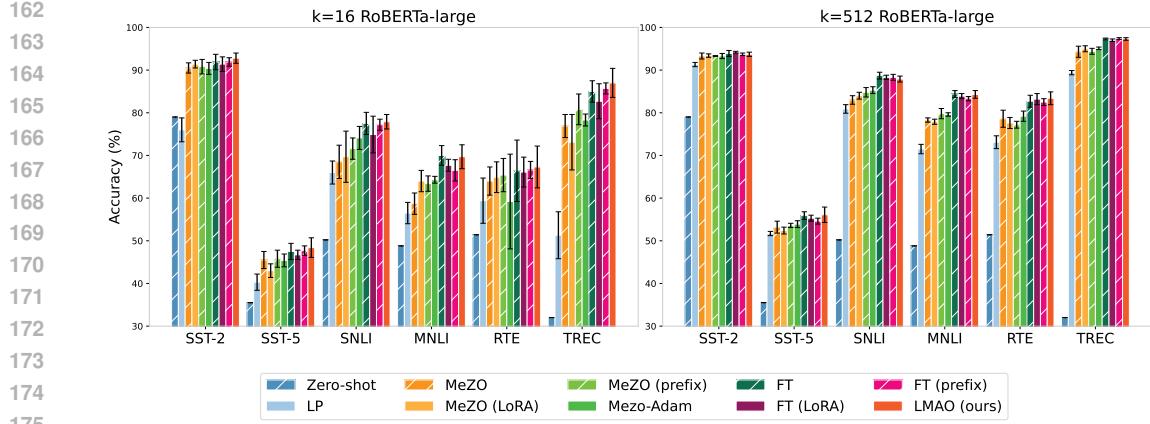


Figure 3: Experiments on RoBERTa-large. We report results for zero-shot, linear probing (LP), MeZO and its variants, our method LMAO, as well as full fine-tuning (FT), LoRA, and prefix-tuning. With $k = 16$, LMAO significantly outperforms LP and MeZO variants. At $k = 512$, LMAO continues to achieve strong performance. Detailed results are shown in Table 1.

these complementary paradigms to enable efficient fine-tuning of LLMs under limited resources while maintaining strong performance.

Each LMAO training iteration consists of two phases: 1) Low-rank adaptation: Low-rank matrices A and B are updated using standard gradient-based optimizers (e.g., AdamW (Loshchilov & Hutter, 2017), SGD), involving a forward pass and a backward propagation to compute exact gradients for the low-rank modules. 2) Zeroth-order optimization: Following the low-rank update, a memory-efficient zeroth-order step adjusts the base model parameters W using two perturbed forward passes to estimate gradients with respect to the global weights.

The full LMAO procedure is summarized in Algorithm 1. Each iteration performs three forward passes—one for the LoRA update and two perturbed passes for zeroth-order gradient estimation—and a single backward pass for LoRA gradients. This design ensures that global weight updates remain full-rank, substantially mitigating the feature loss inherent to purely low-rank or zeroth-order methods.

4 THEORY

In this paper, we consider the following optimization problem:

$$\min_{W, A, B} \mathcal{L}(W; A, B), \quad (1)$$

where $W \in \mathbb{R}^{m \times n}$ represents the model parameters, $A \in \mathbb{R}^{r \times n}$, $B \in \mathbb{R}^{m \times r}$ are low-rank matrices, and \mathcal{L} denotes the loss function.

We briefly introduce the theoretical foundations of the LoRA and MeZO components of our algorithm 1, followed by a theoretical analysis of their alternating optimization.

Firstly, we analyze the theoretical aspects based on the following assumptions.

Assumption 4.1 (Lipschitz smoothness). Let $\mathcal{L}(W; A, B)$ be a differentiable function, there exist constants L_W and L_{BA} such that for any $W_1, W_2 \in \mathbb{R}^{m \times n}$, and $B_1 A_1, B_2 A_2 \in \mathbb{R}^{m \times n}$ the following condition holds:

1. $\|\nabla_W \mathcal{L}(W_1; A, B) - \nabla_W \mathcal{L}(W_2; A, B)\| \leq L_W \|W_1 - W_2\|,$
2. $\|\nabla_{BA} \mathcal{L}(W; A_1, B_1) - \nabla_{BA} \mathcal{L}(W; A_2, B_2)\| \leq L_{BA} \|B_1 A_1 - B_2 A_2\|,$
3. $\|\nabla_W \mathcal{L}(W; A_1, B_1) - \nabla_W \mathcal{L}(W; A_2, B_2)\|_F^2 \leq L_{qua} \|B_1 A_1 - B_2 A_2\|_F^2.$

Assumption 4.2 (Expected smoothness (Khaled & Richtárik, 2020; Malinovsky et al., 2024)). The second moment of the stochastic gradient satisfies

$$\mathbb{E} [\|\nabla \mathcal{L}(W; A, B)\|_F^2] \leq 2\alpha (\mathcal{L}(W; A, B) - \mathcal{L}^*) + \beta \|\nabla \mathcal{L}(W; A, B)\|_F^2$$

Algorithm 1 LMAO (Low-rank and Memory-efficient Zeroth-Order Alternating Optimization)

Input: Model Weight Parameters W , loss \mathcal{L} , step budget T , perturbation scale ε , batch size \mathcal{B} , learning rate scheduler $\{\eta_t\}$.

```

216 1: for  $t = 1, \dots, T$  do
217 2:   Sample batch  $\mathcal{B} \subset \mathcal{D}$  and random seed  $s$ 
218 3:    $[A, B] = \text{LoRA\_Update}(W, A, B)$                                  $\triangleright$  Forward pass and backpropagation
219 4:    $W \leftarrow \text{PerturbParameters}(W, \varepsilon, s)$ 
220 5:    $\mathcal{L}^+ \leftarrow \mathcal{L}(W, A, B; \mathcal{B})$ 
221 6:    $W \leftarrow \text{PerturbParameters}(W, -2\varepsilon, s)$                        $\triangleright$  Forward pass w/ positive perturbation
222 7:    $\mathcal{L}^- \leftarrow \mathcal{L}(W, A, B; \mathcal{B})$                            $\triangleright$  Forward pass w/ negative perturbation
223 8:    $W \leftarrow \text{PerturbParameters}(W, \varepsilon, s)$ 
224 9:   projected_grad  $\leftarrow (\mathcal{L}^+ - \mathcal{L}^-)/(2\varepsilon)$ 
225 10:  Reset random number generator with seed  $s$ 
226 11:  for  $W_i \in W$  do
227 12:     $z \sim \mathcal{N}(0, 1)$ 
228 13:     $W_i \leftarrow W_i - \eta_t \cdot \text{projected\_grad} \cdot z$ 
229 14:  end for
230 15: end for
231 16: function LoRA_Update( $W, A, B$ )
232 17:    $[A, B] = [A, B] - \eta_t \nabla_{AB} \mathcal{L}(W, A, B; \mathcal{B})$ 
233 18:   return  $[A, B]$ 
234 19: end function
235 20: function PerturbParameters( $W, \varepsilon, s$ )
236 21:   Reset random number generator with seed  $s$ 
237 22:   for  $W_i \in W$  do
238 23:      $z \sim \mathcal{N}(0, 1)$ 
239 24:      $W_i \leftarrow W_i + \varepsilon z$ 
240 25:   end for
241 26:   return  $W$ 
242 27: end function

```

246 for some $\alpha, \beta \geq 0$ and all $W \in \mathbb{R}^{m \times n}$, $B \in \mathbb{R}^{m \times r}$ and $A \in \mathbb{R}^{r \times n}$.

248 Part of LoRA. We begin by discussing the LoRA component. The update procedure is as follows:

$$[B_{t+1}, A_{t+1}] = [B_t, A_t] - \eta_{BA} \nabla_{BA} \mathcal{L}(W_t; A_t, B_t), \quad (2)$$

where η is the learning rate, and $\nabla_{BA} \mathcal{L}(W_t; A_t, B_t)$ represents the gradient as expressed below:

$$\nabla_{BA} \mathcal{L}(W_t; A_t, B_t) = [\nabla_B \mathcal{L}(W_t; A_t, B_t), \nabla_A \mathcal{L}(W_t; A_t, B_t)].$$

253 First, under Algorithm 1 and based on Assumption 4.1, we present the descent lemma for LMAO as
254 follows

Lemma 4.3 (Descent lemma). Let \mathcal{L} be a function satisfying Assumption 4.1. Under the condition $\eta < 1/L_{\max}$, where $L_{\max} = \max(L_{BA}, L_W, L_{qua})$. Then we have

$$\begin{aligned} & \mathcal{L}(W_{t+1}; A_{t+1}, B_{t+1}) \leq \mathcal{L}(W_t; A_t, B_t) \\ & - \frac{\eta}{2} \left(\left\| \nabla_{BA} \mathcal{L}(W_t; A_t, B_t) \right\|_F^2 + \left\| \nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1}) \right\|_F^2 \right). \end{aligned}$$

Part of MeZO. The MeZO method is an efficient memory optimization approach that combines zeroth-order methods. Given a labeled dataset \mathcal{D} and a mini-batch \mathcal{M} of size $|\mathcal{M}|$, we introduce the classic zeroth-order gradient as follows

Definition 4.4 (Zeroth Gradient). Given a model with parameters $W \in \mathbb{R}^{m \times n}$ and loss function \mathcal{L} , SPSA approximates the gradient on a minibatch \mathcal{M} as

$$\begin{aligned} \hat{\nabla}_W \mathcal{L}(W; A, B; \mathcal{M}) &= \frac{\mathcal{L}(W + \varepsilon z; A, B; \mathcal{M}) - \mathcal{L}(W - \varepsilon z; A, B; \mathcal{M})}{2\varepsilon} z \\ &\approx zz^T \nabla_W \mathcal{L}(W; A, B; \mathcal{M}). \end{aligned} \quad (3)$$

where $z \in \mathbb{R}^{m \times n}$, $z \sim \mathcal{N}(0, I_{m \times n})$ and $\varepsilon \rightarrow 0^+$ is the perturbation scale.

270 The update procedure of MeZO is as follows:
 271

$$272 \quad W_{t+1} = W_t - \eta_W \hat{\nabla}_W \mathcal{L}(W_t; A_{t+1}, B_{t+1}; \mathcal{M}). \quad (4)$$

273 **Definition 4.5.** The SGD gradient estimate on a minibatch of size \mathcal{M} has covariance $\sum_{\mathcal{M}}$ as
 274

$$275 \quad |\mathcal{M}| \left(\mathbb{E}[\nabla \mathcal{L}(W; A, B; \mathcal{M}) \nabla^T \mathcal{L}(W; A, B; \mathcal{M})] - \nabla \mathcal{L}(W; A, B) \nabla^T \mathcal{L}(W; A, B) \right).$$

278 $\hat{\nabla}_W \mathcal{L}(W; A, B; \mathcal{M})$ be a statistic, we need to examine the unbiasedness of its corresponding estimator,
 279 which leads to the following lemmas.

280 **Lemma 4.6 (Unbiasedness).** Based on Definition 4.4, for a mini-batch \mathcal{M} , we have

$$281 \quad \mathbb{E}_{z \sim \mathcal{N}(0, I)} [\hat{\nabla}_W \mathcal{L}(W; A, B; \mathcal{M})] = \hat{\nabla}_W \mathcal{L}(W; A, B).$$

283 **Lemma 4.7 (Biased Estimate).** Based on Definition 4.4, the gradient norm of MeZO is

$$285 \quad \mathbb{E}_{z \sim \mathcal{N}(0, I)} [\|\hat{\nabla}_W \mathcal{L}(W; A, B; \mathcal{M})\|_F^2] = \frac{mn + N - 1}{N} \mathbb{E}[\|\hat{\nabla}_W \mathcal{L}(W; A, B)\|_F^2],$$

286 where m, n is the size of parameters and N is the number of z sampled in SPSA (Definition 4.4).
 287

288 **Alternating Optimization.** In Algorithm 1, we optimize by alternately applying LoRA and MeZO
 289 methods. To analyze the convergence, we require the following lemma:

291 **Lemma 4.8.** Let $\mathcal{L}(W; A, B)$ be the L -smoothness function, we have

$$292 \quad \|\nabla \mathcal{L}(W_t; A_t, B_t)\|_F^2 \leq 2 \left(1 + \frac{L_{BA}}{L_{\max}} \right)^2 (\|\nabla_{BA} \mathcal{L}(W_t; A_t, B_t)\|_F^2 + \|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2).$$

295 Under Assumptions 4.1 and Assumption 4.2, the convergence of Algorithm 1 is guaranteed by the
 296 following Theorem:

297 **Theorem 4.9.** The iterates of with SGD satisfy

$$299 \quad \min_{0 \leq t \leq T-1} \mathbb{E} [\|\nabla \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2] \leq \frac{6(1 + \frac{L_{BA}}{L_{\max}})^2 (\mathcal{L}(W_0; A_0, B_0) - \mathcal{L}^*)}{\eta_{\max} T},$$

302 where the stepsize $\eta_{\max} := \max(\eta_{BA}, \eta_W)$ satisfy

$$304 \quad \eta_{\max} \leq \min \left(\frac{N}{1000\beta_2 L_W(mn + N - 1)}, \frac{1000N}{N\beta_1 L_{BA} + 2\alpha_2 L_W(mn + N - 1)}, \right. \\ 305 \quad \left. \frac{4}{L_{BA}}, \sqrt{\frac{N}{T(N\alpha_1 L_{BA} + \alpha_2 L_W(mn + N - 1))}} \right)$$

309 and is set according to the experimental configuration as specified in expression $\frac{\eta_{\min}}{\eta_{\max}} \geq \frac{1}{1000}$.
 310

311 Building on the finite series construction technique introduced in (Malinovsky et al., 2024) for the
 312 theoretical analysis of LoRA, we incorporate SGD-based zeroth-order optimization and tangential
 313 scaling inequalities in an alternating scheme (Lu et al., 2019). This leads to an adaptive convergence
 314 analysis, with detailed proofs provided in the appendix.
 315

316 5 EXPERIMENTS

318 In this section, we present comprehensive experiments on multiple fine-tuning tasks across models
 319 with different parameter scales, aiming to rigorously evaluate the performance of LMAO. The
 320 experimental setup is detailed as follows:
 321

322 **Models.** The experiments are evaluated across multiple language models, including the medium-sized
 323 RoBERTa-Large(Liu et al., 2019) and the large autoregressive Open Pre-trained Transformers (OPT)
 324 series(Zhang et al., 2022), comprising OPT-1.3B, OPT-2.7B, and OPT-6.7B.

Datasets. We conduct experiments across multiple fine-tuning tasks in GLUE (Wang et al., 2018), SuperGLUE (Wang et al., 2019) benchmark and other datasets. For RoBERTa-Large, we conduct experiments on SST-2 (Socher et al., 2013), SST-5 (Socher et al., 2013), SNLI (Bowman et al., 2015), MNLI (Williams et al., 2017), RTE (Dagan et al., 2005; Haim et al., 2006; Giampiccolo et al., 2007; Bentivogli et al., 2009), and TREC (Voorhees & Tice, 2000). For OPT-1.3B, we conduct experiments on SST-2 (Socher et al., 2013), CB (De Marneffe et al., 2019), BoolQ (Clark et al., 2019), WSC (Levesque et al., 2012), WIC (Pilehvar & Camacho-Collados, 2018), MultiRC (Khashabi et al., 2018), COPA (Roemmele et al., 2011), and ReCoRD (Zhang et al., 2018).

Baselines. The baselines compared in our experiments contain both gradient-free-based and gradient-based methods, including: 1) Zero-shot, 2) In-context learning, 3) Memory-efficient zeroth-order optimizer (MeZO), 4) Low-rank adaptation (LoRA), and 5) Fine-tuning (Full-parameter). The available code for the baselines is from <https://github.com/princeton-nlp/MeZO>.

5.1 MEDIUM-SIZED MASKED LANGUAGE MODELS

We first conduct experiments using the masked language model RoBERTa-Large. We compare the performance of the LMAO with other methods on sentiment analysis, natural language inference, and topic classification tasks. We adopt the classic approaches from previous studies, including few-shot and many-shot settings. For $k = 16$ and $k = 512$, we sample k cases per class. A uniform fine-tuning of $1K$ steps is applied. The results are summarized in Table 1 and Figure 3.

LMAO significantly outperforms Zero-shot, LP, MeZO, and its variants. Across all six tasks, LMAO consistently optimizes the pre-trained model and outperforms Zero-shot, LP, MeZO, and its various extensions. On several tasks, it even surpasses full fine-tuning (FT), with the remaining performance gaps within 0.5%.

LMAO approaches the performance of full fine-tuning (FT) as the number of samples increases, with at most a 2% gap. For $k = 16$, LMAO consistently outperforms LoRA across all six tasks, and in most cases even surpasses FT. When increasing k , the performance gap between LMAO and FT narrows further, with some tasks showing less than a 1% difference.

5.2 LARGE-SCALE AUTOREGRESSIVE LANGUAGE MODELS

After achieving promising results with RoBERTa-Large, we extend the LMAO approach to the OPT series models, including three scales with 1.3B, 2.7B, and 6.7B parameters. To systematically evaluate the performance of LMAO, we select several tasks from the SuperGLUE (Wang et al., 2019) benchmark, encompassing text classification, multiple-choice, and generation tasks. In the experiments, we compare LMAO with several state-of-the-art methods, including LoRA, MeZO, ICL (In-context Learning), and Zero-shot learning. To assess the performance differences under identical training conditions, all experiments are conducted with consistent settings, with each model fine-tuned for $1K$ iterations. This setup allows for a more accurate reflection of the strengths and weaknesses of each method. The results for OPT-1.3B are summarized in Table 2 and Figure 2, while additional results for larger models are provided in Appendix A.

LMAO combines efficiency with strong generalization capabilities. On the OPT-1.3B model, it achieves superior performance across a range of SuperGLUE tasks, consistently outperforming mainstream fine-tuning and zero-shot approaches, positioning it as a strong new paradigm.

LMAO (ours) achieves the best performance on the majority of tasks. Ranking first on 7 out of 9 benchmarks, including substantial gains on SST-2, CB, BoolQ, COPA, and ReCoRD. On the remaining two tasks, its performance is competitive, highlighting the robustness and consistency of the approach.

LMAO also serves as a strong PEFT approach. While LoRA is a widely adopted PEFT baseline with competitive performance, LMAO consistently outperforms it on 8 out of 9 tasks, establishing itself as a high-performing and promising PEFT paradigm.

LMAO effectively combines the advantages of zeroth-order optimization and low-rank adaptation. Compared to zero-shot, ICL, and the zeroth-order method MeZO, LMAO demonstrates consistently superior performance. While MeZO generally outperforms Zero-shot and ICL, it still

Table 1: Experiments on RoBERTa-large (350M parameters). LP: Linear probing; ZO, ZO (LoRA), and ZO (prefix): memory-efficient ZO-SGD with full-parameter tuning, LoRA, and prefix-tuning respectively; FT: fine-tuning with Adam. All reported numbers are averaged accuracy (standard deviation). Bold numbers indicate the best performance on each task, while underlined numbers denote the second-best.

Task Type	SST-2	SST-5	SNLI	MNLI	RTE	TREC
	— sentiment —		— natural language inference —			– topic –
Zero-shot	79.0	35.5	50.2	48.8	51.4	32.0
Gradient-free methods: $k = 16$						
LP	76.0 (2.8)	40.3 (1.9)	66.0 (2.7)	56.5 (2.5)	59.4 (5.3)	51.3 (5.5)
MeZO	90.5 (1.2)	45.5 (2.0)	68.5 (3.9)	58.7 (2.5)	64.0 (3.3)	76.9 (2.7)
MeZO (LoRA)	91.4 (0.9)	43.0 (1.6)	69.7 (6.0)	64.0 (2.5)	64.9 (3.6)	73.1 (6.5)
MeZO (prefix)	90.8 (1.7)	45.8 (2.0)	71.6 (2.5)	63.4 (1.8)	65.4 (3.9)	80.8 (3.6)
MeZO-Adam	90.4 (1.4)	45.4 (1.5)	74.1 (2.7)	64.3 (0.8)†	59.2 (11.1)†	78.3 (1.4)
Gradient-based methods: $k = 16$						
FT	91.9 (1.8)	47.5 (1.9)	77.5 (2.6)	70.0 (2.3)	66.4 (7.2)	85.0 (2.5)
FT (LoRA)	91.4 (1.7)	46.7 (1.1)	74.9 (4.3)	67.7 (1.4)	66.1 (3.5)	82.7 (4.1)
FT (prefix)	91.9 (1.0)	47.7 (1.1)	77.2 (1.3)	66.5 (2.5)	66.6 (2.0)	85.7 (1.3)
ours: $k = 16$						
LMAO (ours)	92.8 (1.2)	48.4 (2.3)	77.9 (1.7)	<u>69.7 (2.8)</u>	67.3 (4.9)	87.0 (3.4)
Gradient-free methods: $k = 512$						
LP	91.3 (0.5)	51.7 (0.5)	80.9 (1.0)	71.5 (1.1)	73.1 (1.5)	89.4 (0.5)
MeZO	93.3 (0.7)	53.2 (1.4)	83.0 (1.0)	78.3 (0.5)	78.6 (2.0)	94.3 (1.3)
MeZO (LoRA)	93.4 (0.4)	52.4 (0.8)	84.0 (0.8)	77.9 (0.6)	77.6 (1.3)	95.0 (0.7)
MeZO (prefix)	93.3 (0.1)	53.6 (0.5)	84.8 (1.1)	79.8 (1.2)	77.2 (0.8)	94.4 (0.7)
MeZO-Adam	93.3 (0.6)	53.9 (0.8)	85.3 (0.8)	79.6 (0.4)	79.2 (1.2)	95.1 (0.3)
Gradient-based methods: $k = 512$						
FT	93.9 (0.7)	55.9 (0.9)	88.7 (0.8)	84.4 (0.8)	82.7 (1.4)	97.3 (0.2)
FT (LoRA)	94.2 (0.2)	<u>55.3 (0.7)</u>	<u>88.3 (0.5)</u>	83.9 (0.6)	<u>83.2 (1.3)</u>	97.0 (0.3)
FT (prefix)	93.7 (0.3)	54.6 (0.7)	<u>88.3 (0.7)</u>	83.3 (0.5)	82.5 (0.8)	97.4 (0.2)
ours: $k = 512$						
LMAO (ours)	93.7 (0.5)	56.1 (1.8)	87.9 (0.7)	<u>84.3 (0.9)</u>	83.4 (1.5)	97.3 (0.3)

Table 2: Experiments on OPT-1.3B (with 1000 examples). ICL: in-context learning; LP: linear probing; LoRA: Low-Rank Adaptation; LMAO: Low-rank and Memory-efficient Zeroth-Order Alternating Optimization. LMAO achieves the best performance on 8 out of 9 tasks and ranks second only to LoRA on RTE, with a marginal difference. Bold numbers indicate the best performance on each task, while underlined numbers denote the second-best.

Method	Tasks								
	SST-2	RTE	CB	BoolQ	WSC	WIC	MultiRC	COPA	ReCoRD
	classification						– multiple choice –		
Zero-shot	53.6	53.1	39.3	46.5	43.3	57.5	44.7	<u>75.0</u>	71.3
ICL	74.4	52.9	53.6	54.2	<u>54.5</u>	51.8	45.4	70.3	70.5
MeZO	58.1	53.2	46.4	55.6	49.0	57.9	49.0	74.0	71.4
LoRA	<u>92.2</u>	61.1	<u>67.9</u>	<u>61.4</u>	53.2	<u>58.4</u>	<u>56.3</u>	74.7	<u>72.0</u>
LMAO (ours)	92.4	<u>59.4</u>	69.0	62.0	54.8	58.9	56.7	76.7	72.1

lags behind LMAO—for instance, MeZO achieves comparable results on WiC (57.9% vs. 58.9%) and ReCoRD (71.4% vs. 72.1%), but exhibits notable gaps on other tasks.

5.3 ABLATION STUDY

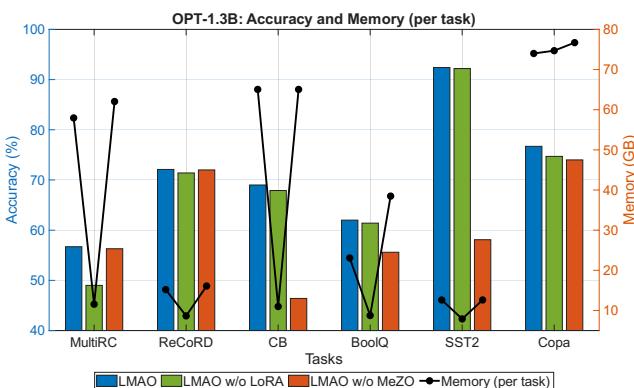


Figure 4: Ablation study results.

components optimized by MeZO in LMAO, and updating only the zeroth-order optimizer components.

We conduct a comprehensive set of ablation studies to systematically evaluate the overall effectiveness of the proposed method. Our analysis compares both accuracy and peak memory consumption across multiple datasets, with all experiments performed on the OPT-1.3B model to ensure consistent and fair comparisons. The ablations cover several key components of the design, including:

- **LMAO w/o LoRA:** Freezing the low-rank LoRA blocks optimized in LMAO, and updating only the zero-order optimizer components.
- **LMAO w/o MeZO:** Freezing the zero-order optimizer

components optimized by MeZO in LMAO, and updating only the low-rank LoRA blocks.

Figure 4 presents the ablation results across several representative tasks. As shown, the full LMAO method consistently achieves substantially higher accuracy than its individual components, namely the low-rank LoRA module and the zeroth-order MeZO module. This indicates that the integration of these two update mechanisms enables LMAO to more effectively capture task-relevant information and enhances the model’s adaptation capability. The performance gains are observed across all evaluated datasets, further validating the effectiveness and robustness of the alternating optimization strategy adopted in LMAO.

Notably, despite combining two distinct update paradigms, LMAO does not incur higher peak memory consumption compared to using LoRA or MeZO alone. This demonstrates that LMAO achieves improved performance without introducing additional memory overhead, maintaining a memory footprint comparable to its component methods. Overall, these findings highlight that LMAO delivers superior task performance while preserving computational efficiency and resource usage, underscoring its practical value in memory-constrained training scenarios.

5.4 PARAMETER SENSITIVITY ANALYSIS

In our method, r and α are two key hyperparameters: r determines the rank of the trainable low-rank matrices, while α scales the contribution of the adapter’s output during fine-tuning. To comprehensively analyze the sensitivity of parameters w.r.t. r and α , we evaluate the RoBERTa-large model on the SST-2 dataset with varying values of r and α . The results are provided in Figure 5. We can see that our method is not sensitive to the settings of r and α . Therefore, in practical applications, we can prioritize smaller values of r to ensure fine-tuning of the model with as few parameters as possible.

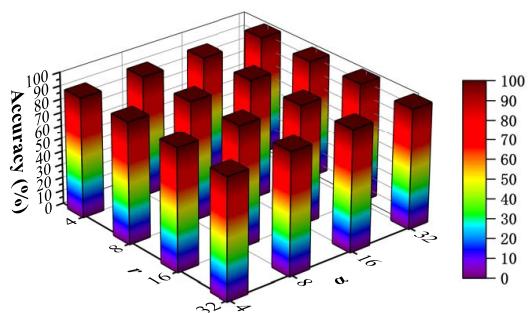


Figure 5: Parameter sensitivity w.r.t. hyperparameters r and α in LMAO.

486 **5.5 MEMORY UTILIZATION ANALYSIS**

488 To examine the memory footprint of
 489 LMAO, we measure the peak GPU memory consumption on a subset of datasets (CB,
 490 SST-2, BoolQ, and COPA) and compare it against two representative baselines: the
 491 low-rank adaptation method LoRA and the zeroth-order optimization method MeZO.
 492 To ensure a fair comparison, all methods are evaluated under identical training con-
 493 figurations and hardware settings. The statistics are summarized in Table 3.

494 As shown in Table 3, the peak memory consumption of LMAO never exceeds that of LoRA or
 495 MeZO, despite integrating both low-rank updates and zeroth-order optimization mechanisms. This
 496 indicates that LMAO does not introduce additional memory burden during training and maintains a
 497 memory footprint comparable to its component methods. Such efficiency is particularly important
 498 in resource-constrained settings, demonstrating that LMAO can be deployed without increasing
 499 hardware requirements.

500
 501 Table 3: Peak GPU Memory(GB) Usage on OPT-
 502 1.3B (1000 Examples). LoRA: Low-Rank Adap-
 503 tation; MeZO: Memory-Efficient Zeroth-Order Opti-
 504 mizer; LMAO: Low-rank and Memory-efficient Zeroth-
 505 Order Alternating Optimization.

Method	SST-2	CB	BoolQ	COPA
MeZO	7.90	10.99	8.75	7.70
LoRA	12.62	65.05	38.48	8.31
LMAO (ours)	12.62	65.05	23.08	8.31

More importantly, LMAO exhibits substantial performance gains across a variety of tasks, consistently outperforming both LoRA and MeZO on classification and multiple-choice benchmarks. These results suggest that LMAO effectively leverages the complementary strengths of the two update strategies, enhancing the model’s learning capability and adaptability while preserving computational efficiency. Overall, the findings strongly support the design philosophy of LMAO, showing that it delivers notable performance improvements

511 while remaining memory-efficient.

514 **6 CONCLUSION**

517 We introduce LMAO, a memory-efficient fine-tuning method for large language models that leverages
 518 full-rank updates. Unlike traditional low-rank adaptation (e.g., LoRA) and zeroth-order optimization
 519 (e.g., MeZO) techniques, which face inherent limitations, LMAO combines the strengths of both
 520 approaches by alternating low-rank matrix updates and zeroth-order forward-pass updates. This
 521 hybrid design significantly enhances model performance while effectively maintaining strict memory
 522 consumption, making it particularly suitable for large language models under resource constraints.
 523 Our theoretical analysis establishes rigorous convergence guarantees, ensuring the robustness and
 524 reliability of the method. Empirical experiments conducted on widely-used models such as RoBERTa-
 525 large and OPT demonstrate that LMAO delivers competitive performance compared to existing
 526 baseline methods. A remaining challenge lies in its training inefficiency on large datasets or long
 527 sequences. We propose that this issue could be addressed through optimization techniques such as
 528 component training, which remains an important avenue for future research and development.

529 **ETHICS STATEMENT**

532 All participants in this work, as well as the paper submission, adhere to the ICLR Code of Ethics (
 533 <https://iclr.cc/public/CodeOfEthics>).

535 **REPRODUCIBILITY STATEMENT**

538 We affirm that the results of this work are fully reproducible. Appendix C provides the theoretical
 539 proofs. Appendix A details the experimental implementations. The source code is available at
<https://anonymous.4open.science/r/mlao-C2EC/>.

540 REFERENCES
541

- 542 Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman,
543 Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report.
544 *arXiv preprint arXiv:2303.08774*, 2023.
- 545 Armen Aghajanyan, Luke Zettlemoyer, and Sonal Gupta. Intrinsic dimensionality explains the
546 effectiveness of language model fine-tuning. *arXiv preprint arXiv:2012.13255*, 2020.
547
- 548 Luisa Bentivogli, Peter Clark, Ido Dagan, and Danilo Giampiccolo. The fifth pascal recognizing
549 textual entailment challenge. *TAC*, 7(8):1, 2009.
- 550
- 551 Samuel R Bowman, Gabor Angeli, Christopher Potts, and Christopher D Manning. A large annotated
552 corpus for learning natural language inference. *arXiv preprint arXiv:1508.05326*, 2015.
- 553 Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal,
554 Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are
555 few-shot learners. *Advances in neural information processing systems*, 33:1877–1901, 2020.
556
- 557 Yiming Chen, Yuan Zhang, Liyuan Cao, Kun Yuan, and Zaiwen Wen. Enhancing zeroth-order
558 fine-tuning for language models with low-rank structures. *arXiv preprint arXiv:2410.07698*, 2024.
- 559
- 560 Christopher Clark, Kenton Lee, Ming-Wei Chang, Tom Kwiatkowski, Michael Collins, and Kristina
561 Toutanova. Boolq: Exploring the surprising difficulty of natural yes/no questions. *arXiv preprint
arXiv:1905.10044*, 2019.
- 562
- 563 Ido Dagan, Oren Glickman, and Bernardo Magnini. The pascal recognising textual entailment
564 challenge. In *Machine learning challenges workshop*, pp. 177–190. Springer, 2005.
565
- 566 Marie-Catherine De Marneffe, Mandy Simons, and Judith Tonhauser. The commitmentbank: In-
567 vestigating projection in naturally occurring discourse. In *proceedings of Sinn und Bedeutung*,
568 volume 23, pp. 107–124, 2019.
- 569
- 570 Tanmay Gautam, Youngsuk Park, Hao Zhou, Parameswaran Raman, and Wooseok Ha. Variance-
571 reduced zeroth-order methods for fine-tuning language models. *arXiv preprint arXiv:2404.08080*,
572 2024.
- 573
- 574 Danilo Giampiccolo, Bernardo Magnini, Ido Dagan, and William B Dolan. The third pascal rec-
575 ognizing textual entailment challenge. In *Proceedings of the ACL-PASCAL workshop on textual
entailment and paraphrasing*, pp. 1–9, 2007.
- 576
- 577 Wentao Guo, Jikai Long, Yimeng Zeng, Zirui Liu, Xinyu Yang, Yide Ran, Jacob R. Gardner, Osbert
578 Bastani, Christopher De Sa, Xiaodong Yu, Beidi Chen, and Zhaozhuo Xu. Zeroth-order fine-
579 tuning of LLMs with extreme sparsity. In *2nd Workshop on Advancing Neural Network Training:
Computational Efficiency, Scalability, and Resource Optimization (WANT@ICML 2024)*, 2024.
580 URL <https://openreview.net/forum?id=pW4MmsnVRq>.
581
- 582 R Bar Haim, Ido Dagan, Bill Dolan, Lisa Ferro, Danilo Giampiccolo, Bernardo Magnini, and Idan
583 Szpektor. The second pascal recognising textual entailment challenge. In *Proceedings of the
Second PASCAL Challenges Workshop on Recognising Textual Entailment*, volume 7, pp. 785–794,
584 2006.
585
- 586
- 587 Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morrone, Quentin De Laroussilhe,
588 Andrea Gesmundo, Mona Attariyan, and Sylvain Gelly. Parameter-efficient transfer learning for
589 nlp. In *International conference on machine learning*, pp. 2790–2799. PMLR, 2019.
- 590
- 591 Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang,
592 Weizhu Chen, et al. Lora: Low-rank adaptation of large language models. *ICLR*, 1(2):3, 2022.
593
- Ahmed Khaled and Peter Richtárik. Better theory for sgd in the nonconvex world. *arXiv preprint
arXiv:2002.03329*, 2020.

- 594 Daniel Khashabi, Snigdha Chaturvedi, Michael Roth, Shyam Upadhyay, and Dan Roth. Looking
 595 beyond the surface: A challenge set for reading comprehension over multiple sentences. In
 596 *Proceedings of the 2018 Conference of the North American Chapter of the Association for Com-*
 597 *putational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pp. 252–262,
 598 2018.
- 599 Brian Lester, Rami Al-Rfou, and Noah Constant. The power of scale for parameter-efficient prompt
 600 tuning. *arXiv preprint arXiv:2104.08691*, 2021.
- 601
- 602 Hector J Levesque, Ernest Davis, and Leora Morgenstern. The winograd schema challenge. *KR*,
 603 2012:13th, 2012.
- 604 Chunyuan Li, Heerad Farkhoor, Rosanne Liu, and Jason Yosinski. Measuring the intrinsic dimension
 605 of objective landscapes. *arXiv preprint arXiv:1804.08838*, 2018.
- 606
- 607 Xiang Lisa Li and Percy Liang. Prefix-tuning: Optimizing continuous prompts for generation. *arXiv*
 608 *preprint arXiv:2101.00190*, 2021.
- 609
- 610 Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike
 611 Lewis, Luke Zettlemoyer, and Veselin Stoyanov. Roberta: A robustly optimized bert pretraining
 612 approach. *arXiv preprint arXiv:1907.11692*, 2019.
- 613
- 614 Yong Liu, Zirui Zhu, Chaoyu Gong, Minhao Cheng, Cho-Jui Hsieh, and Yang You. Sparse mezo: Less
 615 parameters for better performance in zeroth-order llm fine-tuning. *arXiv preprint arXiv:2402.15751*,
 616 2024.
- 617
- 618 Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. *arXiv preprint*
arXiv:1711.05101, 2017.
- 619
- 620 Songtao Lu, Mingyi Hong, and Zhengdao Wang. Pa-gd: On the convergence of perturbed alternating
 621 gradient descent to second-order stationary points for structured nonconvex optimization. In
International Conference on Machine Learning, pp. 4134–4143. PMLR, 2019.
- 622
- 623 Grigory Malinovsky, Umberto Michieli, Hasan Abed Al Kader Hammoud, Taha Ceritli, Hayder
 624 Elesey, Mete Ozay, and Peter Richtárik. Randomized asymmetric chain of lora: The first
 625 meaningful theoretical framework for low-rank adaptation. *arXiv preprint arXiv:2410.08305*,
 2024.
- 626
- 627 Sadhika Malladi, Tianyu Gao, Eshaan Nichani, Alex Damian, Jason D Lee, Danqi Chen, and Sanjeev
 628 Arora. Fine-tuning language models with just forward passes. *Advances in Neural Information*
Processing Systems, 36:53038–53075, 2023.
- 629
- 630 Mohammad Taher Pilehvar and Jose Camacho-Collados. Wic: the word-in-context dataset for
 631 evaluating context-sensitive meaning representations. *arXiv preprint arXiv:1808.09121*, 2018.
- 632
- 633 Melissa Roemmele, Cosmin Adrian Bejan, and Andrew S Gordon. Choice of plausible alterna-
 634 tives: An evaluation of commonsense causal reasoning. In *AAAI spring symposium: logical*
formalizations of commonsense reasoning, pp. 90–95, 2011.
- 635
- 636 Richard Socher, Alex Perelygin, Jean Wu, Jason Chuang, Christopher D Manning, Andrew Y Ng, and
 637 Christopher Potts. Recursive deep models for semantic compositionality over a sentiment treebank.
 638 In *Proceedings of the 2013 conference on empirical methods in natural language processing*, pp.
 1631–1642, 2013.
- 639
- 640 Irene Solaiman, Miles Brundage, Jack Clark, Amanda Askell, Ariel Herbert-Voss, Jeff Wu, Alec
 641 Radford, Gretchen Krueger, Jong Wook Kim, Sarah Kreps, et al. Release strategies and the social
 642 impacts of language models. *arXiv preprint arXiv:1908.09203*, 2019.
- 643
- 644 James C Spall. Multivariate stochastic approximation using a simultaneous perturbation gradient
 645 approximation. *IEEE transactions on automatic control*, 37(3):332–341, 1992.
- 646
- 647 Ellen M Voorhees and Dawn M Tice. Building a question answering test collection. In *Proceedings of*
the 23rd annual international ACM SIGIR conference on Research and development in information
retrieval, pp. 200–207, 2000.

- 648 Alex Wang, Amanpreet Singh, Julian Michael, Felix Hill, Omer Levy, and Samuel R Bowman. Glue:
649 A multi-task benchmark and analysis platform for natural language understanding. *arXiv preprint*
650 *arXiv:1804.07461*, 2018.
- 651
- 652 Alex Wang, Yada Pruksachatkun, Nikita Nangia, Amanpreet Singh, Julian Michael, Felix Hill, Omer
653 Levy, and Samuel Bowman. Superglue: A stickier benchmark for general-purpose language
654 understanding systems. *Advances in neural information processing systems*, 32, 2019.
- 655 Adina Williams, Nikita Nangia, and Samuel R Bowman. A broad-coverage challenge corpus for
656 sentence understanding through inference. *arXiv preprint arXiv:1704.05426*, 2017.
- 657
- 658 Zhiyuan Yu, Yifei Cheng, Liang Ding, Xinmei Tian, Li Shen, and Dacheng Tao. Memory-efficient
659 block coordinate descent for hessian-informed zeroth-order optimizer, 2024a. URL <https://openreview.net/forum?id=q8H9t10Vsy>.
- 660
- 661 Ziming Yu, Pan Zhou, Sike Wang, Jia Li, and Hua Huang. Subzero: Random subspace zeroth-order
662 optimization for memory-efficient llm fine-tuning. *arXiv preprint arXiv:2410.08989*, 2024b.
- 663
- 664 Sheng Zhang, Xiaodong Liu, Jingjing Liu, Jianfeng Gao, Kevin Duh, and Benjamin Van Durme.
665 Record: Bridging the gap between human and machine commonsense reading comprehension.
666 *arXiv preprint arXiv:1810.12885*, 2018.
- 667
- 668 Susan Zhang, Stephen Roller, Naman Goyal, Mikel Artetxe, Moya Chen, Shuhui Chen, Christopher
669 Dewan, Mona Diab, Xian Li, Xi Victoria Lin, et al. Opt: Open pre-trained transformer language
models. *arXiv preprint arXiv:2205.01068*, 2022.
- 670
- 671 Yanjun Zhao, Sizhe Dang, Haishan Ye, Guang Dai, Yi Qian, and Ivor W Tsang. Second-order
672 fine-tuning without pain for llms: A hessian informed zeroth-order optimizer. *arXiv preprint*
arXiv:2402.15173, 2024.
- 673
- 674
- 675
- 676
- 677
- 678
- 679
- 680
- 681
- 682
- 683
- 684
- 685
- 686
- 687
- 688
- 689
- 690
- 691
- 692
- 693
- 694
- 695
- 696
- 697
- 698
- 699
- 700
- 701

APPENDIX

A EXPERIMENT SETUP

A.1 DATASETS

For the RoBERTa-large experiments, we consider a set of standard classification benchmarks, including SST-2 (Socher et al., 2013), SST-5 (Socher et al., 2013), MNLI (Williams et al., 2017), SNLI (Bowman et al., 2015), and RTE (Dagan et al., 2005; Haim et al., 2006; Giampiccolo et al., 2007; Bentivogli et al., 2009). To enable efficient evaluation, we follow the setting in (Malladi et al., 2023) and limit the test set to 1000 examples. For training and validation, we adopt two configurations with $k = 16$ and $k = 512$, where k denotes the number of examples per class for each split.

For the OPT experiments, we evaluate on the SuperGLUE benchmark suite (Wang et al., 2019), which comprises BoolQ (Clark et al., 2019), CB (De Marneffe et al., 2019), COPA (Roemmele et al., 2011), MultiRC (Khashabi et al., 2018), ReCoRD (Zhang et al., 2018), RTE (Dagan et al., 2005; Haim et al., 2006; Giampiccolo et al., 2007; Bentivogli et al., 2009), WIC (Pilehvar & Camacho-Collados, 2018), and WSC (Levesque et al., 2012). In addition, we include SST-2 (Socher et al., 2013). For all tasks, we randomly sample 1000 examples for training, 500 for validation, and 1000 for testing.

A.2 PROMPTS

LMAO adopts the same zeroth-order technique as (Malladi et al., 2023), and therefore we follow their setup for the downstream tasks and prompt templates used to fine-tune RoBERTa-large. The specific templates are provided in Table 4.

Table 4: The prompts used in our RoBERTa-large experiments (see Table 1 and Figure 2) follow the setup in (Malladi et al., 2023), including a template and a set of label words that fill the [MASK] token. The symbols $< S_1 >$ and $< S_2 >$ denote the first and second input sentences, respectively.

Datasets	C	Type	Prompt	Label words
SST-2	2	sentiments cls.	$< S_1 >$ It was [MASK].	{grey, terrible}
SST-5	5	sentiments cls.	$< S_1 >$ It was [MASK].	{grey, good, okey, bad, terrible}
TREC	6	topic cls.	[MASK] : $< S_1 >$.	{Description, Expression, Entity, Human, Location, Number}
MNLI	3	NLI	$< S_1 >$? [MASK], $< S_2 >$	{Yes, Maybe, No}
SNLI	3	NLI	$< S_1 >$? [MASK], $< S_2 >$	{Yes, Maybe, No}
RTE	2	NLI	$< S_1 >$? [MASK], $< S_2 >$	{Yes, No}

Table 5 presents the prompt templates used in our OPT experiments. Specifically, the OPT setup includes two types of tasks: classification and multiple-choice. All prompts follow the setup in (Socher et al., 2013).

A.3 HYPERPARAMETERS

For experiments with the RoBERTa-Large model, we follow the prompt-based fine-tuning paradigm for masked language models as described in (Socher et al., 2013). In OPT experiments, we adopt a similar training strategy to (Socher et al., 2013) for classification tasks, where we extract logits corresponding to the label words and apply cross-entropy loss. For multiple-choice and generative tasks (e.g., QA), we retain only the correct candidate and use teacher forcing on the correct examples. Loss is computed solely on the tokens in the candidate portion, excluding the prompt.

The specific hyperparameter settings for RoBERTa-Large and OPT are provided in Table 6 and Table 7, respectively.

756
 757 Table 5: We use prompt templates in our OPT model experiments, covering classification (cls.) and
 758 multiple-choice (mch.) tasks. Templates follow (Malladi et al., 2023). For multiple-choice, we score
 759 candidates by average log-likelihood and select the highest. For QA, we apply greedy decoding.

760 Dataset	761 Type	762 Prompt
SST-2	cls.	<text>It was terrible/great
RTE	cls.	<preamble> Does this mean that ‘<hypothesis>’ is true? Yes or No? Yes/No/Maybe
CB	cls.	Suppose <premise>Can we infer that ‘<hypothesis>’? Yes, No, or Maybe? Yes/No/Maybe
BoolQ	cls.	<passage><question> Yes/No
WSC	cls.	<text> Does the word ‘<word>’ have the same meaning in these two sentences? Yes, No? Yes/No
WIC	cls.	Does the word ‘<word>’ have the same meaning in these two sentences? Yes, No? <sent1> <sent2> Yes/No
MultiRC	cls.	<paragraph> Question: <question> I found this answer ‘<answer>’. Is that correct? Yes or No? Yes/No
COPA	mch.	<premise>so/because <candidate>
ReCoAD	mch.	<passage> <query>.replace(‘@placeholder’, <candidate>)

784
 785 Table 6: Main hyperparameter settings for all tasks (SST-2, RTE, CB, BoolQ, WSC, WIC, MultiRC,
 786 COPA, ReCoRD) under the RoBERTa-Large model.
 787

788 Method	789 Batch size	790 Learning rate η_{BA}	791 Learning rate η_W	792 Scaling factor α	793 Rank r	794 Perturbation scale ε
LMAO	64	3e-4	1e-7	24	8	1e-3
LoRA	{4, 8, 16}	{1e-4, 3e-4, 5e-4}	/	16	8	/
FT (LoRA)	{4, 8, 16}	{1e-4, 3e-4, 5e-4}	/	16	8	/
MeZo	64	/	{1e-7, 1e-6, 1e-5}	/	/	1e-3
Zero-shot	64	/	/	/	/	/

795 B MORE EXPERIMENT RESULTS

799 B.1 PERFORMANCE ON LARGER MODELS

801 We scale up to the OPT-2.7B and OPT-6.7B model to evaluate LMAO and compare it with several
 802 state-of-the-art methods, including LoRA, MeZo, ICL (in-context learning), and zero-shot learning.
 803 All methods are trained with 1000 iterations under consistent settings. The detailed results are shown
 804 in Table 8 and Table 9.

805 Results show that LMAO performs strongly on OPT-2.7B, achieving superior results on six tasks
 806 (CB, BoolQ, WIC, MultiRC, COPA, ReCoRD). We further scale the model to OPT-6.7B and evaluate
 807 algorithm performance against MeZo and LoRA on a subset of tasks. The results show that LMAO
 808 continues to perform excellently on OPT-6.7B, achieving the best performance on four out of six
 809 tasks (CB, BoolQ, WIC, MultiRC, COPA, ReCoRD). Combined with the results on OPT-2.7B, this
 demonstrates LMAO’s robustness and cross-task generalization, validating its effectiveness and

810
811 Table 7: Hyperparameter settings for all datasets and corresponding methods under the OPT-1.3B
812 model.

813 814 Method	815 816 817 818 Hyperparameters	Tasks								
		SST-2	RTE	CB	BoolQ	819 820 821 WSC	WIC	MultiRC	COPA	ReCoRD
822 823 824 LMAO	Batch size	8	8	8	2	8	8	2	8	2
	Learning rate η_{BA}	1e-5	1e-5	1e-5	1e-5	1e-5	1e-5	1e-5	1e-5	1e-5
	Learning rate η_W	1e-6	1e-6	1e-6	1e-6	1e-6	1e-6	1e-6	1e-6	1e-6
	Scaling factor α	16	16	16	16	16	16	16	16	16
	rank r	8	8	8	8	8	8	8	8	8
825 826 LoRA	Perturbation scale ε	1e-3	1e-3	1e-3	1e-3	1e-3	1e-3	1e-3	1e-3	1e-3
	Batch size	8	8	8	2	8	8	2	8	2
	Learning rate η	1e-5	1e-5	1e-5	1e-5	1e-5	1e-5	1e-5	1e-5	1e-5
	Scaling factor α	16	16	16	16	16	16	16	16	16
827 828 829 MeZO	rank r	8	8	8	8	8	8	8	8	8
	Perturbation scale ε	1e-3	1e-3	1e-3	1e-3	1e-3	1e-3	1e-3	1e-3	1e-3
	Batch size	8	8	8	2	8	8	2	8	2
830 831 832 ICL	Learning rate η	1e-7	1e-7	1e-7	1e-7	1e-7	1e-7	1e-7	1e-7	1e-7
	Perturbation scale ε	1e-3	1e-3	1e-3	1e-3	1e-3	1e-3	1e-3	1e-3	1e-3
	Batch size	32	32	32	32	32	32	32	32	32
833 834 835 Zero-shot	Batch size	64	64	64	64	64	64	64	64	64
	rank r	16	/	16	16	/	/	/	16	/
	Learning rate	8	/	8	8	/	/	/	8	/
836 837 SubZero	Batch size	16	/	16	16	/	/	/	16	/
	Learning rate	1e-5	/	1e-5	1e-5	/	/	/	1e-5	/
	Batch size	16	/	16	16	/	/	/	16	/
838 839 DoRA	Learning rate	1e-5	/	1e-5	1e-5	/	/	/	1e-5	/
	Batch size	16	/	16	16	/	/	/	16	/
	rank r	8	/	8	8	/	/	/	8	/
840 841 842 AdaLoRA	Learning rate	1e-5	/	1e-5	1e-5	/	/	/	1e-5	/
	Batch size	16	/	16	16	/	/	/	16	/
	rank r	8	/	8	8	/	/	/	8	/
	Learning rate	1e-5	/	1e-5	1e-5	/	/	/	1e-5	/

837
838 Table 8: Results on OPT-2.7B (with 1,000 examples). ICL: in-context learning; MeZO: Memory-
839
840 efficient Zeroth-Order optimizer; LoRA: Low-Rank Adaptation; LMAO: Low-rank and Memory-
841 efficient zeroth-order Alternating Optimization. LMAO achieves the best performance across all six
842 tasks. Bold values denote the best result per task.

843 844 Method	Tasks					
	845 846 CB	847 848 BoolQ	849 850 WIC	851 852 MultiRC	COPA	ReCoRD
					853 854 classification	855 856 – multiple choice –
857 858 Zero-shot	50.0	50.9	52.5	43.2	71.0	75.0
859 860 ICL	55.4	56.9	51.6	52.4	77.0	74.0
861 862 MeZO	56.5	59.5	52.4	48.2	73.0	74.9
863 864 LoRA	66.1	68.3	55.5	56.4	83.3	75.2
865 866 LMAO (ours)	68.5	68.5	57.3	57.5	85.3	75.3

adaptability on larger models. Thus, LMAO proves to be an efficient and memory-friendly fine-tuning approach for large language models.

This demonstrates LMAO’s robustness and cross-task generalization, validating its effectiveness and adaptability on larger models. LMAO thus serves as an efficient and memory-friendly fine-tuning approach for large language models.

B.2 COMPARISON WITH MORE METHODS

Table 9: Results on OPT-6.7B (with 1,000 examples). MeZO: Memory-efficient Zeroth-Order optimizer; LoRA: Low-Rank Adaptation; LMAO: Low-rank and Memory-efficient zeroth-order Alternating Optimization. LMAO leads on seven of nine tasks. Bold values denote the best result per task.

Method	Tasks					
	CB	BoolQ	WIC	MultiRC	COPA	ReCoRD
	classification			multiple choice		
MeZO	50.0	44.2	52.7	49.9	80.3	78.1
LoRA	64.3	49.4	52.9	57.0	81.0	79.7
LMAO (ours)	66.1	65.1	50.0	57.2	81.3	77.6

To enable a more rigorous comparison with recent parameter-efficient fine-tuning (PEFT) approaches, we augment our evaluation with additional zeroth-order and LoRA variants. The experimental results are summarized in Table 10. Based on Table 10, LMAO attains the best overall performance across

Table 10: LMAO vs. PEFT baselines on OPT-1.3B (1000 examples). LOZO: Low-rank Zeroth-Order Stochastic Gradient Descent; SubZero: Random Subspace Zeroth-Order Optimization; DoRA: Dynamic Low-Rank Adaptation; AdaLoRA: Adaptive Budget Allocation for Parameter-Efficient Fine-Tuning; LMAO: Low-rank and Memory-efficient Zeroth-Order Alternating Optimization.

Method	SST-2	CB	BoolQ	COPA
LOZO	86.0	72.0	62.4	69.0
SubZero	63.6	43.0	65.8	76.0
AdaLoRA	87.8	71.4	55.7	76.0
DoRA	87.3	69.4	60.6	75.0
LMAO (ours)	92.4	89.0	62.0	76.7

the four benchmark tasks. Except for isolated cases, it outperforms competing methods; even when a baseline holds a single-task advantage, LMAO’s aggregate superiority remains evident, indicating stronger overall effectiveness and robustness.

C CONVERGENCE ANALYSIS

C.1 SYMBOL SPECIFICATION

The notation used in the derivation of the proofs is summarized in Table 11.

C.2 PROOF OF LEMMA 4.3

Proof. Using the total differential formula, we have

$$\nabla \mathcal{L}(W_t; A_{t+1}, B_{t+1}) = [\nabla_{BA} \mathcal{L}(W_t; A_{t+1}, B_{t+1}), \nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})].$$

Table 11: Table of notations.

Variable	Definition
\mathcal{L}	The loss function.
\mathcal{L}^*	The theoretical optimum of the objective function.
A	Low-rank matrix in $\mathbb{R}^{r \times n}$.
B	Low-rank matrix in $\mathbb{R}^{m \times r}$.
W	The weight parameter matrix of a pre-trained model.
η_{BA}	The learning rate of the low-rank module in Algorithm 1.
η_W	The learning rate of the zeroth-order module in Algorithm 1.
L_{BA}, L_W, L_{qua}	The Lipschitz constant.
\mathcal{M}	A minibatch sample in SPSA-based sampling.
$ \mathcal{M} $	The number of samples \mathcal{M} in a minibatch.
$\hat{\nabla}\mathcal{L}$	The update gradient from algorithmic sampling.
m, n	The dimensionality of the pre-trained model's parameter matrix.
N	The number of samples drawn in SPSA.
α, β	The smoothness constant in the expected smoothness assumption.
$\text{Sym}(\cdot)$	The symmetrization function for tensors.
\otimes	Tensor product.

Under Assumption 4.1, we have (descent lemma)

$$\begin{aligned}
& \mathcal{L}(W_{t+1}; A_{t+1}, B_{t+1}) \\
& \leq \mathcal{L}(W_t; A_t, B_t) + \nabla_{BA}^T \mathcal{L}(W_t; A_t, B_t)([B_{t+1}, A_{t+1}] - [B_t, A_t]) \\
& \quad + \frac{L_{BA}}{2} \| [B_{t+1}, A_{t+1}] - [B_t, A_t] \|_F^2 + \nabla_W^T \mathcal{L}(W_t; A_{t+1}, B_{t+1})(W_{t+1} - W_t) \\
& \quad + \frac{L_W}{2} \| W_{t+1} - W_t \|_F^2 \\
& \stackrel{(a)}{\leq} \mathcal{L}(W_t; A_t, B_t) - \eta_{BA} (\| \nabla_{BA} \mathcal{L}(W_t; A_t, B_t) \|_F^2 + \| \nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1}) \|_F^2) \\
& \quad + \frac{\eta_{BA}^2 L_{BA}}{2} \| \nabla_{BA} \mathcal{L}(W_t; A_t, B_t) \|_F^2 + \frac{\eta_{BA}^2 L_W}{2} \| \nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1}) \|_F^2 \\
& \stackrel{(b)}{\leq} \mathcal{L}(W_t; A_t, B_t) - \frac{\eta_{BA}}{2} (\| \nabla_{BA} \mathcal{L}(W_t; A_t, B_t) \|_F^2 + \| \nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1}) \|_F^2),
\end{aligned}$$

where (a) is true because of the update rule of gradient descent in each block and Assumption 4.1, in (b) we used $\eta_{BA} \leq 1/L_{\max}$ and $L_{\max} := \max(L_{BA}, L_W)$. \square

C.3 PROOF OF LEMMA 4.6

Proof. Using the equation equation 3 we note that

$$\begin{aligned}
\mathbb{E}_{z \sim \mathcal{N}(0, I)} [\hat{\nabla}_W \mathcal{L}(W; A, B; \mathcal{M})] &= \mathbb{E}_{z \sim \mathcal{N}(0, I)} [zz^T \nabla_W \mathcal{L}(W; A, B; \mathcal{M})] \\
&= \nabla_W \mathcal{L}(W; A, B; \mathcal{M}) \mathbb{E}_{z \sim \mathcal{N}(0, I)} [zz^T] \\
&= \nabla_W \mathcal{L}(W; A, B; \mathcal{M}) \mathbb{E}_{z \sim \mathcal{N}(0, I)} I \\
&= \nabla_W \mathcal{L}(W; A, B; \mathcal{M}).
\end{aligned}$$

\square

972 C.4 PROOF OF LEMMA 4.7
 973

974 *Proof.* We simplify the notation $\hat{\nabla}_W \mathcal{L}(W_t; A_t, B_t; \mathcal{M}|(x_i, y_i))$ as $\hat{\nabla}_W \mathcal{L}(W_t|(x_i, y_i))$, then we
 975 rewrite the expression as follows:

$$\begin{aligned} & \mathbb{E}_{z \sim \mathcal{N}(0, I)} [\|\hat{\nabla}_W \mathcal{L}(W_t; A_t, B_t; \mathcal{M})\|_F^2] \\ &= \mathbb{E}_{z \sim \mathcal{N}(0, I)} [\hat{\nabla}_W^T \mathcal{L}(W_t; A_t, B_t; \mathcal{M}) \cdot \hat{\nabla}_W \mathcal{L}(W_t; A_t, B_t; \mathcal{M})] \\ &= \frac{1}{|\mathcal{M}|^2 N^2} \sum_{(x_1, y_1), (x_2, y_2) \in \mathcal{M}} \sum_{1 \leq i, j \leq N} \mathbb{E}[\hat{\nabla}_W \mathcal{L}(W_t|(x_1, y_1)) \hat{\nabla}_W^T \mathcal{L}(W_t|(x_2, y_2))] \end{aligned}$$

982 Specifically, if u and v are two independent vectors, then we have
 983

$$\mathbb{E}_{z_i, z_j} [z_i z_i^T u v^T z_j^T z_j] = u v^T,$$

984 when $i \neq j$ and

$$\begin{aligned} \mathbb{E}_{z_i} [z_i z_i^T u v^T z_i^T z_i] &= \mathbb{E}_z [z^{\otimes 4}] \langle u, v \rangle \\ &= \frac{3mn}{mn+2} \mathbf{Sym}(I^{\otimes 2}) \langle u, v \rangle \\ &= \frac{mn}{mn+2} u^T v I + \frac{mn}{mn+2} u v^T. \end{aligned}$$

990 Therefore
 991

$$\begin{aligned} & \mathbb{E}_{z \sim \mathcal{N}(0, I)} [\hat{\nabla}_W^T \mathcal{L}(W_t; A_t, B_t; \mathcal{M}) \cdot \hat{\nabla}_W \mathcal{L}(W_t; A_t, B_t; \mathcal{M})] \\ &= \frac{1}{|\mathcal{M}|^2} \sum_{(x_1, y_1), (x_2, y_2) \in \mathcal{M}} \left(\frac{N-1}{N} + \frac{2mn}{N(mn+2)} \right) \\ &\quad \times \mathbb{E}[(z_i z_i^T \hat{\nabla}_W \mathcal{L}(W_t|(x_1, y_1))) (z_j z_j^T \hat{\nabla}_W^T \mathcal{L}(W_t|(x_2, y_2)))^T] \\ &\quad + \frac{mn}{N(mn+2)} \mathbb{E}[(\hat{\nabla}_W \mathcal{L}(W_t|(x_1, y_1)))^T (\hat{\nabla}_W^T \mathcal{L}(W_t|(x_2, y_2)))] I. \end{aligned}$$

999 Note that when $(x_1, y_1) \neq (x_2, y_2)$, we have
 1000

$$\begin{aligned} & \mathbb{E}[(\hat{\nabla}_W \mathcal{L}(W_t|(x_1, y_1)))^T (\hat{\nabla}_W^T \mathcal{L}(W_t|(x_2, y_2)))] \\ &= \hat{\nabla}_W \mathcal{L}(W_t; A_t, B_t; \mathcal{M}) \hat{\nabla}_W^T \mathcal{L}(W_t; A_t, B_t; \mathcal{M}), \end{aligned}$$

1003 and when $(x_1, y_1) = (x_2, y_2)$, we have
 1004

$$\begin{aligned} & \mathbb{E}[(\hat{\nabla}_W \mathcal{L}(W_t|(x_1, y_1)))^T (\hat{\nabla}_W^T \mathcal{L}(W_t|(x_2, y_2)))] \\ &= \hat{\nabla}_W \mathcal{L}(W_t; A_t, B_t; \mathcal{M}) \hat{\nabla}_W^T \mathcal{L}(W_t; A_t, B_t; \mathcal{M}) + \sum_{\mathcal{M}}, \end{aligned}$$

1008 then we get
 1009

$$\begin{aligned} & \frac{1}{|\mathcal{M}|^2} \sum_{(x_1, y_1), (x_2, y_2) \in \mathcal{M}} \mathbb{E}[(\hat{\nabla}_W \mathcal{L}(W_t|(x_1, y_1)))^T (\hat{\nabla}_W^T \mathcal{L}(W_t|(x_2, y_2)))] \\ &= \nabla_W \mathcal{L}(W_t; A_t, B_t) \nabla_W^T \mathcal{L}(W_t; A_t, B_t) + \sum_{\mathcal{M}}. \end{aligned}$$

1014 Summing these terms yields
 1015

$$\begin{aligned} & \mathbb{E}[(\hat{\nabla}_W \mathcal{L}(W_t; A_t, B_t; \mathcal{M}))^T (\hat{\nabla}_W^T \mathcal{L}(W_t; A_t, B_t; \mathcal{M}))] \\ &= \left(1 + \frac{mn}{N(mn+2)} \right) \left(\hat{\nabla}_W \mathcal{L}(W_t; A_t, B_t) \hat{\nabla}_W^T \mathcal{L}(W_t; A_t, B_t) + \frac{1}{|\mathcal{M}|} \sum_{\mathcal{M}} \right) \\ &\quad + \frac{mn}{N(mn+2)} I \left(\|\hat{\nabla}_W \mathcal{L}(W_t; A_t, B_t)\|_F^2 + \frac{1}{|\mathcal{M}|} \mathbf{tr}(\sum_{\mathcal{M}}) \right). \end{aligned}$$

1022 From above, we obtain
 1023

$$\mathbb{E}[\|\hat{\nabla}_W \mathcal{L}(W_t; A_t, B_t; \mathcal{M})\|_F^2] = \frac{mn+N-1}{N} \mathbb{E}[\|\hat{\nabla}_W \mathcal{L}(W_t; A_t, B_t)\|_F^2].$$

□

1026 C.5 PROOF OF LEMMA 4.8
 1027

1028 *Proof.* By the triangle inequality, we obtain
 1029

$$1030 \|\nabla_W \mathcal{L}(W_t; A_t, B_t)\| \leq \|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1}) - \nabla_W \mathcal{L}(W_t; A_t, B_t)\| \\ 1031 + \|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|.$$

1032 By the Mean inequality, we have
 1033

$$1034 \|\nabla_W \mathcal{L}(W_t; A_t, B_t)\|_F^2 \leq 2\|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1}) - \nabla_W \mathcal{L}(W_t; A_t, B_t)\|_F^2 \\ 1035 + 2\|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2.$$

1037 Combining the updata procedure 2 with Assumption 4.1, we have
 1038

$$1039 \|\nabla_W \mathcal{L}(W_t; A_t, B_t)\|_F^2 \leq 2L_{qua}^2\|[A_{t+1}, B_{t+1}] - [A_t, B_t]\|_F^2 + 2\|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2 \\ 1040 = 2L_{qua}^2\eta_{BA}^2\|\hat{\nabla}_{BA}\mathcal{L}(W_t; A_t, B_t)\|_F^2 + 2\|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2.$$

1041
 1042 For a fixed $0 < \eta_{BA} < \frac{1}{L_{max}}$, where $L_{max} := \max(L_{BA}, L_W, L_{qua})$, the above expression reduces
 1043 to:
 1044

$$1045 \|\nabla_W \mathcal{L}(W_t; A_t, B_t)\|_F^2 \leq 2\left(\frac{L_{qua}^2}{L_{max}^2}\|\hat{\nabla}_{BA}\mathcal{L}(W_t; A_t, B_t)\|_F^2 + \|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2\right).$$

1046 Fixed $L_{max}^2 \geq 2L_{qua}^2$, then the gradient norm of Algorithm 1 satisfies
 1047

$$1048 \|\nabla \mathcal{L}(W_t; A_t, B_t)\|_F^2 = \|\nabla_W \mathcal{L}(W_t; A_t, B_t)\|_F^2 + \|\nabla_{BA}\mathcal{L}(W_t; A_t, B_t)\|_F^2 \\ 1049 \leq 2\left(\frac{L_{qua}^2}{L_{max}^2}\|\nabla_{BA}\mathcal{L}(W_t; A_t, B_t)\|_F^2 + \|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2\right) \\ 1050 + \|\nabla_{BA}\mathcal{L}(W_t; A_t, B_t)\|_F^2 \\ 1051 \leq \left(1 + 2\frac{L_{qua}^2}{L_{max}^2}\right)\|\nabla_{BA}\mathcal{L}(W_t; A_t, B_t)\|_F^2 + 2\|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2 \\ 1052 \leq \left(1 + 2\frac{L_{qua}^2}{L_{max}^2}\right)(\|\nabla_{BA}\mathcal{L}(W_t; A_t, B_t)\|_F^2 + \|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2) \\ 1053 \leq 2\left(1 + \frac{L_{qua}}{L_{max}}\right)^2(\|\nabla_{BA}\mathcal{L}(W_t; A_t, B_t)\|_F^2 + \|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2). \\ 1054 \\ 1055 \\ 1056 \\ 1057 \\ 1058 \\ 1059 \\ 1060 \\ 1061 \\ 1062 \\ 1063$$

1064 \square
 1065

1066 C.6 PROOF OF THEOREM 4.9
 1067

1068 *Proof.* We start from L-smoothness and use equation 2 and equation 4:

$$1069 \mathcal{L}(W_{t+1}; A_{t+1}, B_{t+1}) \leq \mathcal{L}(W_t; A_t, B_t) + \langle \nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1}), W_{t+1} - W_t \rangle \\ 1070 + \langle \nabla_{BA}\mathcal{L}(W_t; A_t, B_t), ([B_{t+1}, A_{t+1}] - [B_t, A_t]) \rangle \\ 1071 + \frac{L_W}{2}(\|[B_{t+1}, A_{t+1}] - [B_t, A_t]\|_F^2) + \frac{L_W}{2}\|W_{t+1} - W_t\|_F^2 \\ 1072 \\ 1073 \\ 1074 \leq \mathcal{L}(W_t; A_t, B_t) - \eta_W \langle \nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1}), \hat{\nabla}_W \mathcal{L}(W_t; A_{t+1}, B_{t+1}) \rangle \\ 1075 - \eta_{BA} \langle \nabla_{BA}\mathcal{L}(W_t; A_t, B_t), \nabla_{BA}\mathcal{L}(W_t; A_t, B_t) \rangle \\ 1076 + \frac{L_{BA}}{2}\eta_{BA}^2\|\nabla_{BA}\mathcal{L}(W_t; A_t, B_t)\|_F^2 \\ 1077 + \frac{L_W}{2}\eta_W^2\|\hat{\nabla}_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2. \\ 1078 \\ 1079$$

1080 Taking conditional expectation and combining Lemma 4.6, Lemma 4.7, we have
1081
$$\begin{aligned} & \mathbb{E}[\mathcal{L}(W_{t+1}; A_{t+1}, B_{t+1})|(W_t; A_t, B_t)] \\ & \leq \mathcal{L}(W_t; A_t, B_t) - \eta_W \|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2 - \eta_{BA} \|\nabla_{BA} \mathcal{L}(W_t; A_t, B_t)\|_F^2 \\ & \quad + \frac{L_{BA}}{2} \eta_{BA}^2 \mathbb{E}[\|\nabla_{BA} \mathcal{L}(W_t; A_t, B_t)\|_F^2] + \frac{L_W}{2} \eta_W^2 \mathbb{E}[\|\hat{\nabla}_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2] \\ & \leq \mathcal{L}(W_t; A_t, B_t) - \eta_{\min} (\|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2 + \|\nabla_{BA} \mathcal{L}(W_t; A_t, B_t)\|_F^2) \\ & \quad + \frac{L_{BA}}{2} \eta_{\max}^2 \mathbb{E}[\|\nabla_{BA} \mathcal{L}(W_t; A_t, B_t)\|_F^2] \\ & \quad + \frac{L_W}{2} \eta_{\max}^2 \frac{mn+N-1}{N} \mathbb{E}[\|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2]. \end{aligned}$$

1087 Subtracting \mathcal{L}^* from both sides of the inequality and using Assumption 4.2, we have

1088
$$\begin{aligned} & \mathbb{E}[\mathcal{L}(W_{t+1}; A_{t+1}, B_{t+1})|(W_t; A_t, B_t)] - \mathcal{L}^* \\ & \leq \mathcal{L}(W_t; A_t, B_t) - \mathcal{L}^* - \eta_{\min} (\|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2 + \|\nabla_{BA} \mathcal{L}(W_t; A_t, B_t)\|_F^2) \\ & \quad + \frac{L_{BA}}{2} \eta_{\max}^2 [2\alpha_1(\mathcal{L}(W_t; A_t, B_t) - \mathcal{L}^*) + \beta_1 \|\nabla_{BA} \mathcal{L}(W_t; A_t, B_t)\|_F^2] \\ & \quad + \frac{L_W}{2} \eta_{\max}^2 \frac{mn+N-1}{N} [2\alpha_2(\mathcal{L}(W_t; A_{t+1}, B_{t+1}) - \mathcal{L}^*) + \beta_2 \|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2]. \end{aligned}$$

1089 Using the Assumption 4.1, we have

1090
$$\begin{aligned} & \mathcal{L}(W_t; A_{t+1}, B_{t+1}) - \mathcal{L}^* = \mathcal{L}(W_t; A_{t+1}, B_{t+1}) - \mathcal{L}(W_t; A_t, B_t) + \mathcal{L}(W_t; A_t, B_t) - \mathcal{L}^* \\ & \leq \langle \nabla_{BA} \mathcal{L}(W_t; A_t, B_t), ([B_{t+1}, A_{t+1}] - [B_t, A_t]) \rangle \\ & \quad + \frac{L_{BA}}{2} \| [B_{t+1}, A_{t+1}] - [B_t, A_t] \|_F^2 + \mathcal{L}(W_t; A_t, B_t) - \mathcal{L}^*. \end{aligned}$$

1091 Define $\delta_{t+1} = \mathbb{E}[\mathcal{L}(W_{t+1}; A_{t+1}, B_{t+1})|[B_t, A_t]] - \mathcal{L}^*$ and combining the above equations, then we have

1092
$$\begin{aligned} & \delta_{t+1} \leq \delta_t - \eta_{\min} (\|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2 + \|\nabla_{BA} \mathcal{L}(W_t; A_t, B_t)\|_F^2) \\ & \quad + \alpha_1 L_{BA} \eta_{\max}^2 \delta_t + \frac{1}{2} L_{BA} \eta_{\max}^2 \beta_1 \|\nabla_{BA} \mathcal{L}(W_t; A_t, B_t)\|_F^2 \\ & \quad + \alpha_2 L_W \eta_{\max}^2 \frac{mn+N-1}{N} (\mathcal{L}(W_t; A_{t+1}, B_{t+1}) - \mathcal{L}(W_t; A_t, B_t)) \\ & \quad + \alpha_2 L_W \eta_{\max}^2 \frac{mn+N-1}{N} \delta_t + \frac{1}{2} \beta_2 L_W \eta_{\max}^2 \|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2 \\ & \leq \left(1 + \alpha_1 L_{BA} \eta_{\max}^2 + \alpha_2 L_W \eta_{\max}^2 \frac{mn+N-1}{N} \right) \delta_t \\ & \quad - \eta_{\max} \left(\frac{\eta_{\min}}{\eta_{\max}} - \frac{1}{2} \beta_2 L_W \eta_{\max}^2 \frac{mn+N-1}{N} \right) \|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2 \\ & \quad - \eta_{\max} \left(\frac{\eta_{\min}}{\eta_{\max}} - \frac{1}{2} L_{BA} \eta_{\max}^2 \beta_1 - \alpha_2 L_W \eta_{\max}^2 \frac{mn+N-1}{N} \left(-\eta_{\min} + \frac{L_{BA}}{2} \eta_{\max}^2 \right) \right) \\ & \quad \times \|\nabla_{BA} \mathcal{L}(W_t; A_t, B_t)\|_F^2. \end{aligned}$$

1093 Combining the above equations, we have

1094
$$\begin{aligned} & \frac{1}{2000} \eta_{\max} (\|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2 + \|\nabla_{BA} \mathcal{L}(W_t; A_t, B_t)\|_F^2) \\ & \leq \left(1 + \alpha_1 L_{BA} \eta_{\max}^2 + \alpha_2 L_W \eta_{\max}^2 \frac{mn+N-1}{N} \right) \delta_t - \delta_{t+1}, \end{aligned}$$

1095 where η satisfies

1096
$$\begin{cases} \frac{\eta_{\min}}{\eta_{\max}} - \frac{1}{2} \beta_2 L_W \eta_{\max}^2 \frac{mn+N-1}{N} \geq \frac{1}{2000} \\ \frac{\eta_{\min}}{\eta_{\max}} - \frac{1}{2} L_{BA} \eta_{\max}^2 \beta_1 - \alpha_2 L_W \eta_{\max}^2 \left(-\eta_{\min} + \frac{L_{BA}}{2} \eta_{\max}^2 \right) \geq \frac{1}{2000} \\ -\eta_{\min} + \frac{L_{BA}}{2} \eta_{\max}^2 \leq \eta_{\max} \end{cases}$$

1134 i.e.

1135

$$\eta_{\max} \leq \min \left(\frac{N}{1000\beta_2 L_W(mn + N - 1)}, \frac{1000N}{N\beta_1 L_{BA} + 2\alpha_2 L_W(mn + N - 1)}, \frac{4}{L_{BA}} \right).$$

1136 Let fix $p_{-1} > 0$ and define $p_t = p_{t-1} \left(1 + \alpha_1 L_{BA} \eta_{\max}^2 + \alpha_2 L_W \eta_{\max}^2 \frac{mn + N - 1}{N} \right)^{-1}$ for all
 1137
 1138
 1139
 1140
 1141
 1142
 $t \geq 0$. Multiplying by $\frac{p_t}{\eta_{\max}}$, we have

1143

$$\frac{p_t}{2} (\|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2 + \|\nabla_{BA} \mathcal{L}(W_t; A_t, B_t)\|_F^2) \leq (p_{t-1} \delta_t - p_t \delta_{t+1}).$$

1144
 1145 Summing up both sides as $t = 0, 1, 2, \dots, T - 1$, we have

1146

$$\sum_{t=0}^{T-1} \frac{p_t}{2} (\|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2 + \|\nabla_{BA} \mathcal{L}(W_t; A_t, B_t)\|_F^2) \\ \leq (p_{t-1} \delta_t - p_t \delta_{t+1}) \leq p_{-1} \delta_0.$$

1147 Define $P_T = \sum_{t=0}^{T-1} p_t$ and we have

1148

$$P_T = \sum_{t=0}^{T-1} p_t \geq \sum_{t=0}^{T-1} \min\{p_0, p_1, \dots, p_{T-1}\} = T p_{T-1} \\ = \frac{T p_{-1}}{\left(1 + \alpha_1 L_{BA} \eta_{\max}^2 + \alpha_2 L_W \eta_{\max}^2 \frac{mn + N - 1}{N} \right)^T}.$$

1149 Dividing both sides by P_T we have

1150

$$\min_{0 \leq t \leq T-1} \mathbb{E} [\|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2 + \|\nabla_{BA} \mathcal{L}(W_t; A_t, B_t)\|_F^2] \\ \leq \frac{1}{P_T} \sum_{t=0}^{T-1} p_t (\|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2 + \|\nabla_{BA} \mathcal{L}(W_t; A_t, B_t)\|_F^2) \\ \leq \frac{p_{-1} \delta_0}{\eta_{\max} P_T}.$$

1151 Using the fact that $1 + x \leq \exp(x)$, we have

1152

$$\left(1 + \alpha_1 L_{BA} \eta_{\max}^2 + \alpha_2 L_W \eta_{\max}^2 \frac{mn + N - 1}{N} \right)^T \\ \leq \exp T \left(\alpha_1 L_{BA} \eta_{\max}^2 + \alpha_2 L_W \eta_{\max}^2 \frac{mn + N - 1}{N} \right) \\ \leq \exp(1) \leq 3,$$

1153 where the second inequality holds because $\eta_{\max} \leq \sqrt{N / (T(N\alpha_1 L_{BA} + \alpha_2 L_W(mn + N - 1)))}$
 1154 by assumption. Combining the above equations, we have

1155

$$\min_{0 \leq t \leq T-1} \mathbb{E} [\|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2 + \|\nabla_{BA} \mathcal{L}(W_t; A_t, B_t)\|_F^2] \leq \frac{3(\mathcal{L}(W_0; A_0, B_0) - \mathcal{L}^*)}{\eta_{\max} T}.$$

1156 By the Lemma 4.8, we have

1157

$$\|\nabla \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2 \\ \leq 2 \left(1 + \frac{L_{BA}}{L_{\max}} \right)^2 (\|\nabla_{BA} \mathcal{L}(W_t; A_t, B_t)\|_F^2 + \|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2).$$

1188 Taking expectations and minimizing over t on both sides of the inequality, then we have
 1189

$$\begin{aligned} 1190 \min_{0 \leq t \leq T-1} & \mathbb{E} \left[\|\nabla \mathcal{L}(W_t; A_t, B_t)\|_F^2 \right] \\ 1191 & \leq \min_{0 \leq t \leq T-1} 2 \left(1 + \frac{L_{BA}}{L_{\max}} \right)^2 \left(\mathbb{E} \left[\|\nabla_{BA} \mathcal{L}(W_t; A_t, B_t)\|_F^2 \right] + \mathbb{E} \left[\|\nabla_W \mathcal{L}(W_t; A_{t+1}, B_{t+1})\|_F^2 \right] \right) \\ 1192 & \leq \frac{6(1 + \frac{L_{BA}}{L_{\max}})^2 (\mathcal{L}(W_0; A_0, B_0) - \mathcal{L}^*)}{\eta_{\max} T}. \\ 1193 & \\ 1194 & \\ 1195 & \\ 1196 & \\ 1197 & \end{aligned}$$

□

1200 D USE OF LLMs

1201 LLMs are used for grammar checking and language polishing.
 1202

1203
 1204
 1205
 1206
 1207
 1208
 1209
 1210
 1211
 1212
 1213
 1214
 1215
 1216
 1217
 1218
 1219
 1220
 1221
 1222
 1223
 1224
 1225
 1226
 1227
 1228
 1229
 1230
 1231
 1232
 1233
 1234
 1235
 1236
 1237
 1238
 1239
 1240
 1241